

# lstmproject

May 11, 2025

## 1 Libraries

```
[18]: import os
import numpy as np
import pandas as pd
from tqdm import tqdm
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split

import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("device:", device)
```

device: cuda

## 2 Data Loading

```
[19]: # from google.colab import drive
# drive.mount('/content/drive')

data_path = "/kaggle/input/car-data/car_data" # "/content/drive/MyDrive/
↪car_data"

[20]: def load(data_path, max_files=None):
    X_all, y_all = [], []

    files = sorted([f for f in os.listdir(data_path) if f.endswith('.csv')])
    if max_files:
        files = files[:max_files]

    for fname in tqdm(files):
        df = pd.read_csv(os.path.join(data_path, fname))
        if df.shape[0] < 67:
```

```

        continue

    # Separate scalers for input (12 features) and target (2 features)
    input_scaler = MinMaxScaler()
    target_scaler = MinMaxScaler()

    input_seq = input_scaler.fit_transform(df.iloc[:62].values) # 62×12
    target_seq = target_scaler.fit_transform(df.iloc[62:67][['Local_X', 'Local_Y']].values) # 5×2

    X_all.append(input_seq)
    y_all.append(target_seq)

return np.array(X_all), np.array(y_all)

```

### 3 Data Splitting

```

[21]: X_data, y_data = load(data_path, max_files=9400)

X_train, X_temp, y_train, y_temp = train_test_split(X_data, y_data, test_size=0.3, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)

print(f"Train: {X_train.shape}, {y_train.shape}")
print(f"Val: {X_val.shape}, {y_val.shape}")
print(f"Test: {X_test.shape}, {y_test.shape}")

```

100% | 9400/9400 [00:38<00:00, 247.17it/s]

Train: (6580, 62, 12), (6580, 5, 2)

Val: (1410, 62, 12), (1410, 5, 2)

Test: (1410, 62, 12), (1410, 5, 2)

### 4 Dataset Loaders

```

[22]: class CarDataset(Dataset):
    def __init__(self, X, y):
        self.X = torch.tensor(X, dtype=torch.float32)
        self.y = torch.tensor(y, dtype=torch.float32)

    def __len__(self):
        return len(self.X)

    def __getitem__(self, idx):
        return self.X[idx], self.y[idx]

```

```

batch_size = 128
train_loader = DataLoader(CarDataset(X_train, y_train), batch_size=batch_size,
    ↪shuffle=True)
val_loader = DataLoader(CarDataset(X_val, y_val), batch_size=batch_size)
test_loader = DataLoader(CarDataset(X_test, y_test), batch_size=batch_size)

```

## 5 LSTM Custom Design

```

[23]: class LSTM(nn.Module):
    def __init__(self, input_size=12, hidden_size=128, output_size=2,
    ↪output_steps=5, dropout=0.2):
        super().__init__()
        self.input_size = input_size
        self.hidden_size = hidden_size
        self.output_size = output_size
        self.output_steps = output_steps
        self.dropout = nn.Dropout(dropout)
        self.projection = nn.Linear(hidden_size, output_size)

        # LSTM Gate weights and biases
        self.W_xi = nn.Parameter(torch.empty(input_size, hidden_size))
        self.W_hi = nn.Parameter(torch.empty(hidden_size, hidden_size))
        self.b_i = nn.Parameter(torch.zeros(hidden_size))

        self.W_xf = nn.Parameter(torch.empty(input_size, hidden_size))
        self.W_hf = nn.Parameter(torch.empty(hidden_size, hidden_size))
        self.b_f = nn.Parameter(torch.zeros(hidden_size))

        self.W_xc = nn.Parameter(torch.empty(input_size, hidden_size))
        self.W_hc = nn.Parameter(torch.empty(hidden_size, hidden_size))
        self.b_c = nn.Parameter(torch.zeros(hidden_size))

        self.W_xo = nn.Parameter(torch.empty(input_size, hidden_size))
        self.W_ho = nn.Parameter(torch.empty(hidden_size, hidden_size))
        self.b_o = nn.Parameter(torch.zeros(hidden_size))

        # Xavier
        for w in [self.W_xi, self.W_hi, self.W_xf, self.W_hf,
                  self.W_xc, self.W_hc, self.W_xo, self.W_ho]:
            nn.init.xavier_uniform_(w)

    def lstm_cell(self, x_t, h_prev, c_prev):
        i = torch.sigmoid(x_t @ self.W_xi + h_prev @ self.W_hi + self.b_i)
        f = torch.sigmoid(x_t @ self.W_xf + h_prev @ self.W_hf + self.b_f)
        g = torch.tanh(x_t @ self.W_xc + h_prev @ self.W_hc + self.b_c)

```

```

o = torch.sigmoid(x_t @ self.W_xo + h_prev @ self.W_ho + self.b_o)

c_t = f * c_prev + i * g
h_t = o * torch.tanh(c_t)
return self.dropout(h_t), c_t

def forward(self, x):
    batch_size, seq_len, _ = x.size()
    h_t = torch.zeros(batch_size, self.hidden_size, device=x.device)
    c_t = torch.zeros(batch_size, self.hidden_size, device=x.device)

    for t in range(seq_len):
        h_t, c_t = self.lstm_cell(x[:, t, :], h_t, c_t)

    # predicting next 5 datapoints
    outputs = []
    x_next = x[:, -1, :]
    for _ in range(self.output_steps):
        h_t, c_t = self.lstm_cell(x_next, h_t, c_t)
        pred = self.projection(h_t)
        outputs.append(pred)
        x_next = torch.cat([pred, x_next[:, 2:]], dim=1)

    return torch.stack(outputs, dim=1) # (batch, 5, 2)

```

## 6 Training, Validation and Testing Loops

```

[24]: def RMSELoss(pred, target):
    return torch.sqrt(torch.mean((pred - target) ** 2))

# Training
def train_model(config, model, train_loader, val_loader, device='cuda'):
    lr = config["lr"]
    epochs = config["epochs"]
    name = config["name"]

    model.to(device)
    optimizer = torch.optim.Adam(model.parameters(), lr=lr)
    scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer,
mode='min', factor=0.5, patience=10)
    criterion = RMSELoss

    train_rmse = []
    val_rmse = []

    for epoch in range(1, epochs + 1):

```

```

model.train()
running_train_loss = 0

for xb, yb in train_loader:
    xb, yb = xb.to(device), yb.to(device)
    pred = model(xb)
    loss = criterion(pred, yb)

    optimizer.zero_grad()
    loss.backward()
    torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
    optimizer.step()
    running_train_loss += loss.item()

avg_train_loss = running_train_loss / len(train_loader)
train_rmse.append(avg_train_loss)

# Validation
model.eval()
running_val_loss = 0
with torch.no_grad():
    for xb, yb in val_loader:
        xb, yb = xb.to(device), yb.to(device)
        pred = model(xb)
        loss = criterion(pred, yb)
        running_val_loss += loss.item()
avg_val_loss = running_val_loss / len(val_loader)
val_rmse.append(avg_val_loss)

scheduler.step(avg_val_loss)

if epoch % 10 == 0:
    print(f"Epoch {epoch}/{epochs} | {name} | Train RMSE:␣
↪{avg_train_loss:.4f} | Val RMSE: {avg_val_loss:.4f} | LR: {optimizer.
↪param_groups[0]['lr']:.6f}")

    return model, train_rmse, val_rmse

# Testing
def test(model, test_loader, device='cuda'):
    model.eval()
    total_loss = 0
    with torch.no_grad():
        for xb, yb in test_loader:
            xb, yb = xb.to(device), yb.to(device)
            pred = model(xb)
            loss = RMSELoss(pred, yb)

```

```

        total_loss += loss.item()
    avg_loss = total_loss / len(test_loader)
    return avg_loss

```

## 7 Experiments and Results

```

[25]: def experiments(configs):
    results = []

    for cfg in configs:
        print(f"\n config: {cfg['name']}")
        model = LSTM(
            input_size=12,
            hidden_size=cfg["hidden_size"],
            output_size=2,
            output_steps=5,
            dropout=cfg["dropout"]
        )

        model, train_loss, val_loss = train_model(cfg, model, train_loader,
        ↪ val_loader, device=device)
        test_loss = test(model, test_loader, device=device)

        results.append({
            "name": cfg["name"],
            "train_loss": train_loss,
            "val_loss": val_loss,
            "test_loss": test_loss
        })

        plt.figure(figsize=(10, 4))
        epochs = range(1, len(train_loss) + 1)
        plt.plot(epochs, train_loss, 'b--', label='Train RMSE')
        plt.plot(epochs, val_loss, 'r-', label='Val RMSE')
        plt.title(f"RMSE Curve: {cfg['name']}")
        plt.xlabel("Epoch")
        plt.ylabel("RMSE")
        plt.legend()
        plt.grid(True)
        plt.tight_layout()
        plt.show()

    print("\n Final Test RMSEs:")
    for r in results:
        print(f"{r['name']:<25} - Test RMSE: {r['test_loss']:.4f}")


```

```
return results
```

## 8 Testing on Different Model Configurations

```
[28]: configs = [  
    {"name": "model_1", "hidden_size": 128, "lr": 1e-3, "epochs": 200, "dropout":  
    ↪ 0.2},  
    {"name": "model_2", "hidden_size": 256, "lr": 1e-3, "epochs": 100, "dropout":  
    ↪ 0.2},  
    {"name": "model_3", "hidden_size": 512, "lr": 1e-3, "epochs": 100, "dropout":  
    ↪ 0.2},  
    {"name": "model_4", "hidden_size": 512, "lr": 1e-3, "epochs": 80, "dropout":  
    ↪ 0.3},  
    {"name": "model_5", "hidden_size": 1024, "lr": 1e-3, "epochs": 100,  
    ↪ "dropout": 0.2},  
    {"name": "model_6", "hidden_size": 1024, "lr": 1e-3, "epochs": 60, "dropout":  
    ↪ 0.4},  
]  
results = experiments(configs)
```

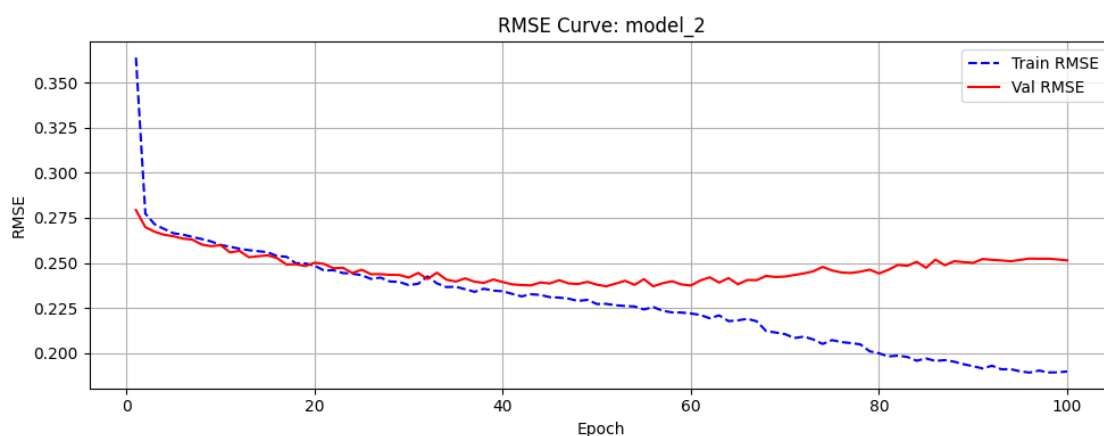
```
config: model_1  
Epoch 10/200 | model_1 | Train RMSE: 0.2638 | Val RMSE: 0.2609 | LR: 0.001000  
Epoch 20/200 | model_1 | Train RMSE: 0.2526 | Val RMSE: 0.2500 | LR: 0.001000  
Epoch 30/200 | model_1 | Train RMSE: 0.2451 | Val RMSE: 0.2464 | LR: 0.001000  
Epoch 40/200 | model_1 | Train RMSE: 0.2401 | Val RMSE: 0.2457 | LR: 0.001000  
Epoch 50/200 | model_1 | Train RMSE: 0.2350 | Val RMSE: 0.2426 | LR: 0.001000  
Epoch 60/200 | model_1 | Train RMSE: 0.2329 | Val RMSE: 0.2417 | LR: 0.001000  
Epoch 70/200 | model_1 | Train RMSE: 0.2264 | Val RMSE: 0.2424 | LR: 0.000500  
Epoch 80/200 | model_1 | Train RMSE: 0.2218 | Val RMSE: 0.2426 | LR: 0.000250  
Epoch 90/200 | model_1 | Train RMSE: 0.2195 | Val RMSE: 0.2442 | LR: 0.000125  
Epoch 100/200 | model_1 | Train RMSE: 0.2175 | Val RMSE: 0.2444 | LR: 0.000125  
Epoch 110/200 | model_1 | Train RMSE: 0.2169 | Val RMSE: 0.2439 | LR: 0.000063  
Epoch 120/200 | model_1 | Train RMSE: 0.2172 | Val RMSE: 0.2440 | LR: 0.000031  
Epoch 130/200 | model_1 | Train RMSE: 0.2156 | Val RMSE: 0.2438 | LR: 0.000016  
Epoch 140/200 | model_1 | Train RMSE: 0.2156 | Val RMSE: 0.2440 | LR: 0.000008  
Epoch 150/200 | model_1 | Train RMSE: 0.2146 | Val RMSE: 0.2440 | LR: 0.000004  
Epoch 160/200 | model_1 | Train RMSE: 0.2153 | Val RMSE: 0.2441 | LR: 0.000002  
Epoch 170/200 | model_1 | Train RMSE: 0.2160 | Val RMSE: 0.2441 | LR: 0.000001  
Epoch 180/200 | model_1 | Train RMSE: 0.2148 | Val RMSE: 0.2441 | LR: 0.000000  
Epoch 190/200 | model_1 | Train RMSE: 0.2147 | Val RMSE: 0.2441 | LR: 0.000000  
Epoch 200/200 | model_1 | Train RMSE: 0.2153 | Val RMSE: 0.2441 | LR: 0.000000
```



```

config: model_2
Epoch 10/100 | model_2 | Train RMSE: 0.2600 | Val RMSE: 0.2599 | LR: 0.001000
Epoch 20/100 | model_2 | Train RMSE: 0.2485 | Val RMSE: 0.2501 | LR: 0.001000
Epoch 30/100 | model_2 | Train RMSE: 0.2377 | Val RMSE: 0.2419 | LR: 0.001000
Epoch 40/100 | model_2 | Train RMSE: 0.2343 | Val RMSE: 0.2394 | LR: 0.001000
Epoch 50/100 | model_2 | Train RMSE: 0.2273 | Val RMSE: 0.2380 | LR: 0.001000
Epoch 60/100 | model_2 | Train RMSE: 0.2220 | Val RMSE: 0.2375 | LR: 0.001000
Epoch 70/100 | model_2 | Train RMSE: 0.2106 | Val RMSE: 0.2425 | LR: 0.000500
Epoch 80/100 | model_2 | Train RMSE: 0.1999 | Val RMSE: 0.2442 | LR: 0.000250
Epoch 90/100 | model_2 | Train RMSE: 0.1928 | Val RMSE: 0.2501 | LR: 0.000125
Epoch 100/100 | model_2 | Train RMSE: 0.1899 | Val RMSE: 0.2514 | LR: 0.000063

```




```

config: model_3
Epoch 10/100 | model_3 | Train RMSE: 0.2578 | Val RMSE: 0.2578 | LR: 0.001000
Epoch 20/100 | model_3 | Train RMSE: 0.2453 | Val RMSE: 0.2522 | LR: 0.001000

```




Epoch 30/100 | model\_3 | Train RMSE: 0.2374 | Val RMSE: 0.2423 | LR: 0.001000  
Epoch 40/100 | model\_3 | Train RMSE: 0.2301 | Val RMSE: 0.2355 | LR: 0.001000  
Epoch 50/100 | model\_3 | Train RMSE: 0.2214 | Val RMSE: 0.2338 | LR: 0.001000  
Epoch 60/100 | model\_3 | Train RMSE: 0.2091 | Val RMSE: 0.2407 | LR: 0.001000  
Epoch 70/100 | model\_3 | Train RMSE: 0.1864 | Val RMSE: 0.2471 | LR: 0.000500  
Epoch 80/100 | model\_3 | Train RMSE: 0.1685 | Val RMSE: 0.2519 | LR: 0.000250  
Epoch 90/100 | model\_3 | Train RMSE: 0.1568 | Val RMSE: 0.2560 | LR: 0.000125  
Epoch 100/100 | model\_3 | Train RMSE: 0.1525 | Val RMSE: 0.2598 | LR: 0.000063




config: model\_4

Epoch 10/80 | model\_4 | Train RMSE: 0.2595 | Val RMSE: 0.2567 | LR: 0.001000  
Epoch 20/80 | model\_4 | Train RMSE: 0.2511 | Val RMSE: 0.2471 | LR: 0.001000  
Epoch 30/80 | model\_4 | Train RMSE: 0.2416 | Val RMSE: 0.2452 | LR: 0.001000  
Epoch 40/80 | model\_4 | Train RMSE: 0.2352 | Val RMSE: 0.2388 | LR: 0.001000  
Epoch 50/80 | model\_4 | Train RMSE: 0.2298 | Val RMSE: 0.2343 | LR: 0.001000  
Epoch 60/80 | model\_4 | Train RMSE: 0.2244 | Val RMSE: 0.2369 | LR: 0.001000  
Epoch 70/80 | model\_4 | Train RMSE: 0.2124 | Val RMSE: 0.2383 | LR: 0.000500  
Epoch 80/80 | model\_4 | Train RMSE: 0.2011 | Val RMSE: 0.2428 | LR: 0.000250




config: model\_5

Epoch 10/100	model_5	Train RMSE: 0.2580	Val RMSE: 0.2575	LR: 0.001000
Epoch 20/100	model_5	Train RMSE: 0.2438	Val RMSE: 0.2508	LR: 0.001000
Epoch 30/100	model_5	Train RMSE: 0.2360	Val RMSE: 0.2396	LR: 0.001000
Epoch 40/100	model_5	Train RMSE: 0.2274	Val RMSE: 0.2392	LR: 0.001000
Epoch 50/100	model_5	Train RMSE: 0.2192	Val RMSE: 0.2372	LR: 0.001000
Epoch 60/100	model_5	Train RMSE: 0.2044	Val RMSE: 0.2370	LR: 0.001000
Epoch 70/100	model_5	Train RMSE: 0.1779	Val RMSE: 0.2419	LR: 0.000500
Epoch 80/100	model_5	Train RMSE: 0.1535	Val RMSE: 0.2562	LR: 0.000250
Epoch 90/100	model_5	Train RMSE: 0.1420	Val RMSE: 0.2586	LR: 0.000125
Epoch 100/100	model_5	Train RMSE: 0.1355	Val RMSE: 0.2605	LR: 0.000063



config: model\_6

Epoch 10/60	model_6	Train RMSE: 0.2641	Val RMSE: 0.2686	LR: 0.001000
Epoch 20/60	model_6	Train RMSE: 0.2556	Val RMSE: 0.2556	LR: 0.001000
Epoch 30/60	model_6	Train RMSE: 0.2455	Val RMSE: 0.2459	LR: 0.001000
Epoch 40/60	model_6	Train RMSE: 0.2389	Val RMSE: 0.2408	LR: 0.001000
Epoch 50/60	model_6	Train RMSE: 0.2327	Val RMSE: 0.2366	LR: 0.001000
Epoch 60/60	model_6	Train RMSE: 0.2276	Val RMSE: 0.2365	LR: 0.001000



Final Test RMSEs:

model_1	- Test RMSE: 0.2498
model_2	- Test RMSE: 0.2598
model_3	- Test RMSE: 0.2689
model_4	- Test RMSE: 0.2462
model_5	- Test RMSE: 0.2707
model_6	- Test RMSE: 0.2375